

Polynomial Regression Implementation

Documentation Generated by Karan Sehgal

Table of Contents

1. Introduction
2. Mathematical Theory
3. Implementation Overview
4. PolynomialRegression Class
5. Workflow Function
6. Usage Examples
7. Visualizations
8. Conclusion

1. Introduction

This document provides comprehensive documentation for the Polynomial Regression implementation from scratch. The implementation includes:

- Complete polynomial regression algorithm without relying on machine learning libraries
- Feature transformation and model training using normal equations
- Comprehensive evaluation metrics (MSE, MAE, R^2)
- Visualization capabilities for model analysis
- Automated documentation generation

Key Features:

- Support for arbitrary polynomial degrees
- Detailed model evaluation metrics
- Comparison of multiple polynomial degrees
- Residual analysis
- Professional PDF documentation generation

The implementation follows software engineering best practices including type hints, proper error handling, and modular design.

2. Mathematical Theory

Polynomial Regression extends linear regression by modeling the relationship between the independent variable (x) and dependent variable (y) as an nth degree polynomial. The general form of the polynomial regression model is:

$$y = \beta_0 + \beta_1x + \beta_2x^2 + \dots + \beta_nx^n + \varepsilon$$

Where: • y: Dependent variable (target) • x: Independent variable (feature) • $\beta_0, \beta_1, \dots, \beta_n$: Model coefficients • ε : Error term **Model Training:** The coefficients are calculated using the normal equation:

$$\beta = (X^T X)^{-1} X^T y$$

3. Implementation Overview

The implementation consists of two main components: **1. PolynomialRegression Class** • Core polynomial regression functionality • Feature transformation • Model training via normal equations • Prediction and evaluation **2. polynomial_regression_workflow Function** • Complete analysis workflow • Handles data splitting • Trains multiple models with different degrees • Generates comprehensive visualizations • Produces evaluation metrics The implementation uses only NumPy for core mathematical operations and matplotlib for visualization.

4. PolynomialRegression Class

The PolynomialRegression class implements the core polynomial regression algorithm.

Class Definition:

```
class PolynomialRegression:
    """
    Polynomial Regression implementation from scratch.

    Attributes:
        degree (int): Degree of the polynomial
        coefficients (np.ndarray): Model coefficients after fitting
    """

    def __init__(self, degree: int = 2) -> None:
        """Initialize with polynomial degree"""
        self.degree = degree
        self.coefficients = None
```

Key Methods:

fit(X, y)

Trains the model using normal equations

Parameter	Description
X	Input features (1D or 2D array)
y	Target values

Returns: self (for method chaining)

predict(X)

Makes predictions using fitted model

Parameter	Description
X	Input features to predict on

Returns: Array of predicted values

score(X, y)

Calculates R² score

Parameter	Description
X	Input features
y	True target values

Returns: R² score (float)

5. Workflow Function

The `polynomial_regression_workflow` function provides a complete analysis pipeline:

Function Definition:

```
def polynomial_regression_workflow(
    X: np.ndarray,
    y: np.ndarray,
    degrees: List[int] = [1, 2, 3, 5],
    test_size: float = 0.2,
    random_state: int = 42,
    true_function: Optional[Callable] = None,
    title: str = "Polynomial Regression Analysis"
) -> Dict:
    """
    Complete polynomial regression analysis workflow.

    Returns dictionary containing:
    - best_model: Best performing model
    - best_degree: Degree of best model
    - all_models: All trained models
    - metrics: Evaluation metrics
    - figures: Generated matplotlib figures
    """
```

Workflow Steps:

Step	Description
1. Data Preparation	Splits data into training/test sets
2. Model Training	Trains models for each polynomial degree
3. Evaluation	Calculates MSE, MAE, R^2 for each model
4. Visualization	Generates comparison plots and residual analysis
5. Model Selection	Selects best model based on test performance

6. Usage Examples

Basic Usage:

```
# Create and fit model
model = PolynomialRegression(degree=3)
model.fit(X_train, y_train)

# Make predictions
predictions = model.predict(X_test)

# Evaluate model
r2 = model.score(X_test, y_test)
```

Complete Workflow Example:

```
# Generate synthetic data
np.random.seed(42)
X = np.linspace(-3, 3, 100)
y = 2 + 3*X - 1.5*X**2 + 0.5*X**3 + np.random.normal(0, 2, 100)

# Run complete analysis
results = polynomial_regression_workflow(
    X, y,
    degrees=[1, 2, 3, 4, 5],
    true_function=lambda x: 2 + 3*x - 1.5*x**2 + 0.5*x**3
)

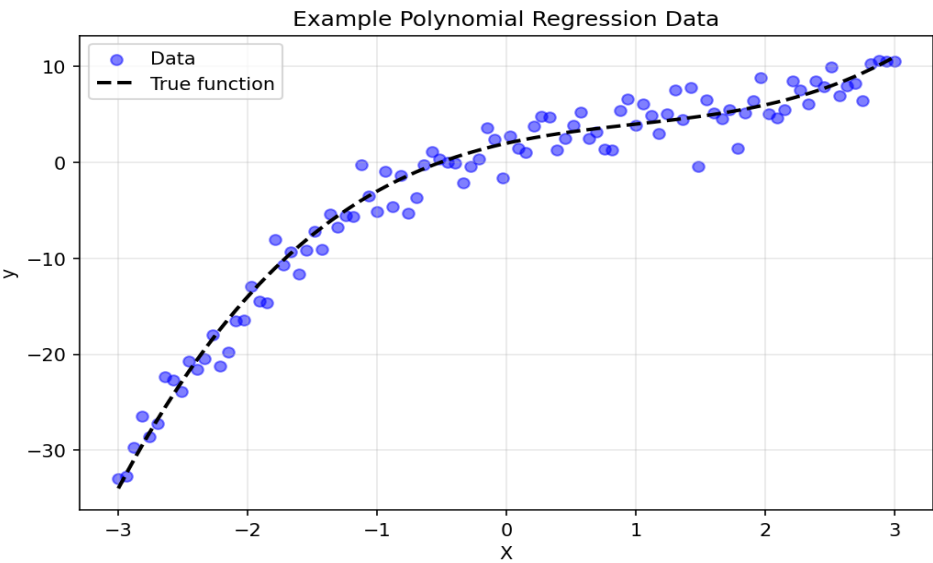
# Access results
print(f"Best degree: {results['best_degree']}")
print(f"Best R² score: {results['metrics'][results['best_degree']]['test']['R²']:.3f}")
```

7. Visualizations

The workflow generates several visualizations to analyze model performance:

Visualization	Description
Model Comparison	Shows fits of different polynomial degrees
Metrics Comparison	Compares R^2 and MSE across degrees
Residual Analysis	Plots residuals for best model

Example Visualization:



8. Conclusion

This implementation provides a complete polynomial regression solution with:

- **Mathematical Rigor:** Proper implementation of normal equations
- **Comprehensive Analysis:** Multiple evaluation metrics and visualizations
- **Professional Documentation:** Self-explanatory PDF documentation

The modular design makes it easy to extend for specific use cases while maintaining the core mathematical correctness.